

Exploratory Data Analysis (EDA)

Justin Post

Big Picture

- Big Data characteristics
- Course split into four topics
 1. Programming in `python`
 2. Big Data Management
 3. Modeling Big Data (with `Spark` via `pyspark`)
 4. Streaming Data

Uses for Data

Four major goals with data:

1. Description
2. Prediction/Classification
3. Inference
4. Pattern Finding

- First step of most analysis is to get to know your data! Done through an **Exploratory Data Analysis**

EDA

- Essentially **Descriptive Statistics** with a bit more big picture stuff about your data
- EDA generally consists of a few steps:
 - Understand how your data is stored
 - Do basic data validation
 - Determine rate of missing values
 - Clean data up data as needed
 - Investigate distributions
 - Univariate measures/graphs
 - Multivariate measures/graphs
 - Apply transformations and repeat previous step

Understand how your data is stored

- Should know if your data has read in how you think it should!

Read in some data (we'll learn more about this later!)

```
import pandas as pd
wine_data = pd.read_csv("https://www4.stat.ncsu.edu/~online/datasets/winequality-full.csv")
wine_data.head()

##   fixed acidity  volatile acidity  citric acid ... alcohol  quality type
## 0            7.4                 0.70      0.00 ...     9.4      5  Red
## 1            7.8                 0.88      0.00 ...     9.8      5  Red
## 2            7.8                 0.76      0.04 ...     9.8      5  Red
## 3           11.2                 0.28      0.56 ...     9.8      6  Red
## 4            7.4                 0.70      0.00 ...     9.4      5  Red
##
## [5 rows x 13 columns]
```

Understand how your data is stored

- Should know if your data has read in how you think it should!

```
wine_data.info()

## <class 'pandas.core.frame.DataFrame'>
## RangeIndex: 6497 entries, 0 to 6496
## Data columns (total 13 columns):
##   #   Column           Non-Null Count  Dtype  
##   --  --  
##   0   fixed acidity    6497 non-null   float64 
##   1   volatile acidity 6497 non-null   float64 
##   2   citric acid      6497 non-null   float64 
##   3   residual sugar   6497 non-null   float64 
##   4   chlorides        6497 non-null   float64 
##   5   free sulfur dioxide 6497 non-null   float64 
##   6   total sulfur dioxide 6497 non-null   float64 
##   7   density          6497 non-null   float64 
##   8   pH               6497 non-null   float64 
##   9   sulphates        6497 non-null   float64 
##   10  alcohol          6497 non-null   float64 
##   11  quality          6497 non-null   int64  
##   12  type              6497 non-null   object 
## dtypes: float64(11), int64(1), object(1)
## memory usage: 660.0+ KB
```

Do basic data validation

- Usually look at quick summary stats of all the data to check that things make sense

```
wine_data.describe()
```

```
##      fixed acidity  volatile acidity   ...    alcohol    quality
##  count      6497.00000      6497.00000   ...    6497.00000  6497.00000
##  mean       7.215307       0.339666   ...    10.491801   5.818378
##  std        1.296434       0.164636   ...    1.192712   0.873255
##  min        3.800000       0.080000   ...     8.000000   3.000000
##  25%        6.400000       0.230000   ...    9.500000   5.000000
##  50%        7.000000       0.290000   ...   10.300000   6.000000
##  75%        7.700000       0.400000   ...   11.300000   6.000000
##  max        15.900000      1.580000   ...   14.900000   9.000000
## 
## [8 rows x 12 columns]
```

Determine rate of missing values

- Every programming language has indicators for missing values
- In python, we use `NaN` for 'not a number' (in `pandas`) (might use other things for missing with other data objects/modules)

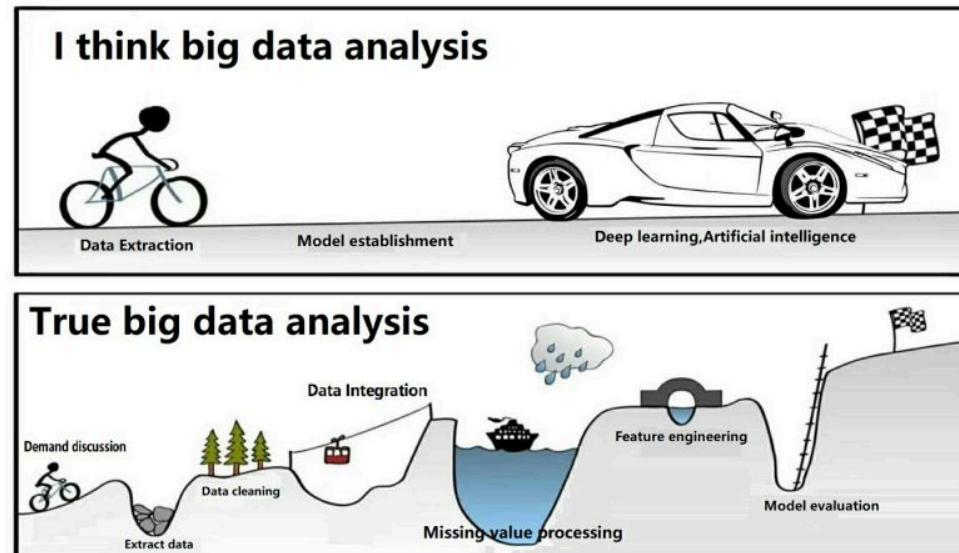
```
wine_data.isnull().sum()
```

```
## fixed acidity      0
## volatile acidity   0
## citric acid        0
## residual sugar     0
## chlorides          0
## free sulfur dioxide 0
## total sulfur dioxide 0
## density            0
## pH                 0
## sulphates          0
## alcohol            0
## quality            0
## type               0
## dtype: int64
```

Clean data up data as needed

May need to

- reread data with different specifications
- fill missing values
- remove some rows and/or columns
- check your data against some gold standard?



Investigate distributions

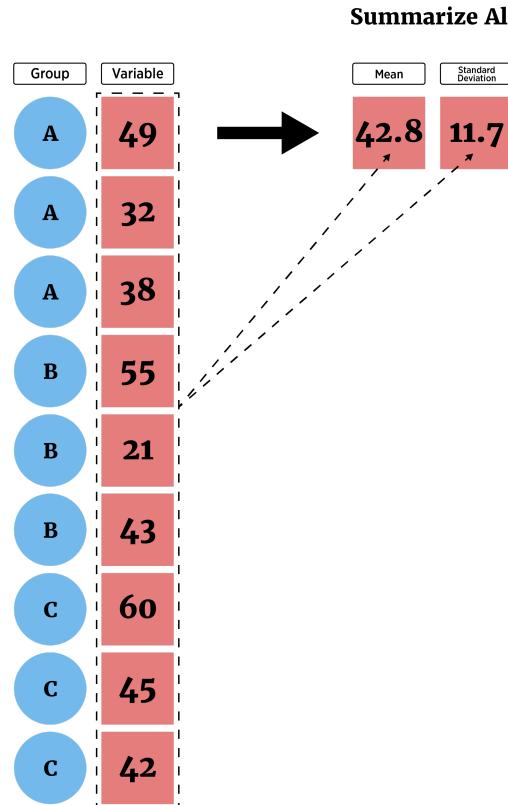
Goal: Understand types of data and their distributions

- Univariate measures/graphs
- Multivariate measures/graphs

Investigate distributions

Goal: Understand types of data and their distributions

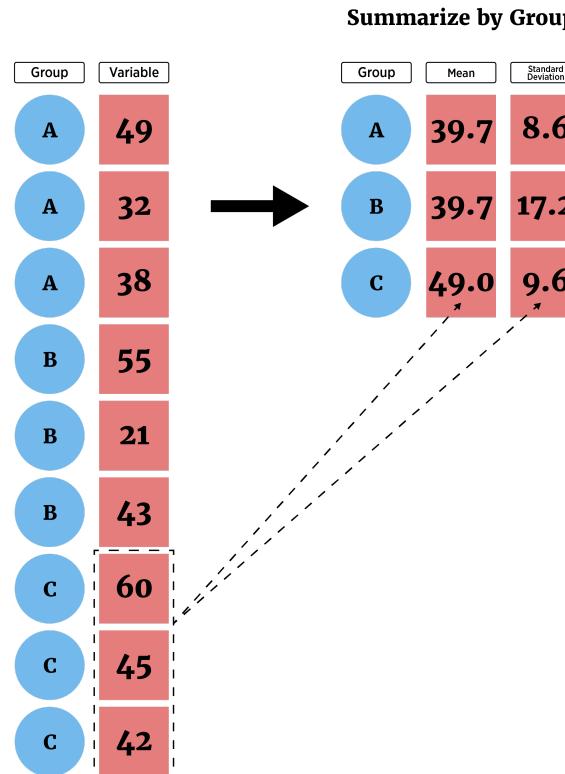
- Numerical summaries



Making Sense of Data

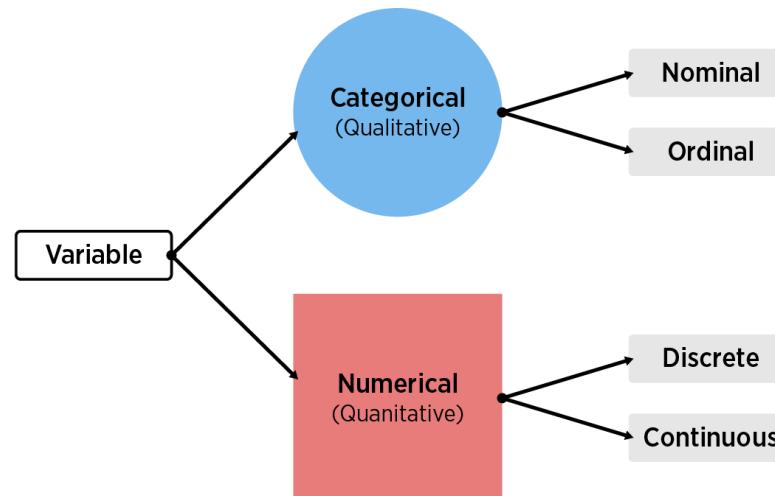
Goal: Understand types of data and their distributions

- Numerical summaries (across subgroups)



Types of Data

- How to summarize data depends on the type of data
 - Categorical (Qualitative) variable - entries are a label or attribute
 - Numeric (Quantitative) variable - entries are a numerical value where math can be performed



Making Sense of Data

Goal: Understand types of data and their distributions

- Numerical summaries (across subgroups)
 - Contingency Tables
 - Mean/Median
 - Standard Deviation/Variance/IQR
 - Quantiles/Percentiles

Categorical Data

Goal: Describe the **distribution** of the variable

- Distribution = pattern and frequency with which you observe a variable
- Categorical variable - entries are a label or attribute
 - Describe the relative frequency (or count) for each category
 - Called a **contingency table**

Categorical Variable Summary - One-way Table

- Count the # of times each category of **one** variable appears!

```
wine_data.type #treat like a numpy array  
  
## 0      Red  
## 1      Red  
## 2      Red  
## 3      Red  
## 4      Red  
##       ...  
## 6492    White  
## 6493    White  
## 6494    White  
## 6495    White  
## 6496    White  
## Name: type, Length: 6497, dtype: object
```

```
sum(wine_data.type == "Red")  
  
## 1599  
  
sum(wine_data.type == "White")  
  
## 4898
```

Categorical Variable Summary - Two-way Table

- Count the # of times each **combination** of categories for *two* variables appear!
- Consider `quality` and `type`

```
sum((wine_data.type == "Red") & (wine_data.quality == 3))  
## 10  
  
sum((wine_data.type == "Red") & (wine_data.quality == 4))  
## 53  
  
sum((wine_data.type == "Red") & (wine_data.quality == 5))  
## 681  
  
#etc
```

Numeric Data

Goal: Describe the **distribution** of the variable

- Distribution = pattern and frequency with which you observe a variable
- Numeric variable - entries are a numerical value where math can be performed

For a single numeric variable, describe the distribution via

- Shape: Histogram, Density plot, ...
- Measures of center: Mean, Median, ...
- Measures of spread: Variance, Standard Deviation, Quartiles, IQR, ...

For two numeric variables, describe the distribution via

- Shape: Scatter plot; Measures of linear relationship: Covariance, Correlation

Numerical Variable Location Summary - Mean

- Sample mean: for a variable in our data set (call it y)

$$\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$$

```
sum(wine_data.alcohol)/len(wine_data.alcohol)  
## 10.491800831152855
```

Numerical Variable Location Summary - Median

- Sample median
 - Sort values
 - Value with 50% of data below and above is the median
 - If even number of observations, average middle two values

```
sorted_alcohol = wine_data.alcohol.sort_values()  
sorted_alcohol  
  
## 4864    8.00  
## 4224    8.00  
## 5438    8.40  
## 5434    8.40  
## 544     8.40  
##  
##      ...  
## 4544   14.00  
## 588    14.00  
## 6102   14.05  
## 5517   14.20  
## 652    14.90  
## Name: alcohol, Length: 6497, dtype: float64
```

```
len(sorted_alcohol)/2  
## 3248.5  
  
sorted_alcohol.values[3248]  
## 10.3
```

Numerical Variable Spread Summary - Variance

- Sample variance is *almost* the average squared deviation from the mean

$$S^2 = \frac{1}{n-1} \sum_{i=1}^n (y_i - \bar{y})^2$$

```
sub = wine_data[0:4].chlorides
sub

## 0    0.076
## 1    0.098
## 2    0.092
## 3    0.075
## Name: chlorides, dtype: float64

mean_chlorides = sum(sub)/4
mean_chlorides

## 0.08525
```

```
sub-mean_chlorides

## 0    -0.00925
## 1    0.01275
## 2    0.00675
## 3   -0.01025
## Name: chlorides, dtype: float64

(sub-mean_chlorides)**2

## 0    0.000086
## 1    0.000163
## 2    0.000046
## 3    0.000105
## Name: chlorides, dtype: float64

sum((sub-mean_chlorides)**2)/3

## 0.000132916666666674
```

Numerical Variable Spread Summary - Standard Deviation

- Sample Standard Deviation = square root of sample variance
 - Puts metric on the scale of the variable

```
import numpy as np
np.sqrt(sum((sub-mean_chlorides)**2)/3)

## 0.011528949070347511
```

Numerical Variable Spread Summary - Quantiles/Percentiles

- Sample quantile - a generalization of the median
 - p^{th} quantile - value with p% of the values below it
 - Also called the $100 \cdot p\%$ ile

```
len(sorted_alcohol)/2  
## 3248.5  
#obtain 0.25 quantile (median of lower half of the data)  
(sorted_alcohol.values[1624]+sorted_alcohol.values[1623])/2  
## 9.5
```

Numerical Variable Relationship Summary - Correlation

- Sample correlation - a measure of the **linear** relationship between two variables
 - Call the variables x and y
 - (x_i, y_i) are numeric variables observed on the same n units, $i = 1, \dots, n$
 - Pearson's correlation coefficient:

$$r = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2 \sum_{i=1}^n (y_i - \bar{y})^2}}$$

Numerical Variable Relationship Summary - Correlation

- Sample correlation - a measure of the **linear** relationship between two variables

```
wine_data.loc[0:4, ["fixed acidity", "chlorides"]]
```

```
##      fixed acidity  chlorides
## 0          7.4     0.076
## 1          7.8     0.098
## 2          7.8     0.092
## 3         11.2     0.075
## 4          7.4     0.076
```

```
wine_data.loc[0:4, ["fixed acidity", "chlorides"]].corr()
```

```
##                  fixed acidity  chlorides
## fixed acidity      1.000000 -0.322814
## chlorides           -0.322814  1.000000
```

Numerical Variable Relationship Summary - Correlation

- Sample correlation - a measure of the **linear** relationship between two variables
 - Sensitive to outliers
 - Spearman's correlation coefficient simply uses Pearson's correlation on the ranks of the data!

```
wine_data.loc[0:4, ["fixed acidity", "chlorides"]]  
##      fixed acidity  chlorides  
## 0          7.4     0.076  
## 1          7.8     0.098  
## 2          7.8     0.092  
## 3         11.2     0.075  
## 4          7.4     0.076
```

Recap

- EDA generally consists of a few steps:
 - Understand how your data is stored
 - Do basic data validation
 - Determine rate of missing values
 - Clean data up data as needed
 - Investigate distributions
 - Univariate measures/graphs
 - Multivariate measures/graphs
 - Apply transformations and repeat previous step
- Usually want summaries for different **subgroups of data!!**