

# LASSO Models

Justin Post

# Recap

- Judge the model's effectiveness using a **Loss** function
- Often split data into a training and test set
  - Perhaps 70/30 or 80/20
- Cross-validation gives a way to use more of the data while still seeing how the model does on test data
  - Commonly 5 fold or 10 fold is done
  - Once a best model is chosen, model is refit on entire data set

# Recap

- Judge the model's effectiveness using a **Loss** function
- Often split data into a training and test set
  - Perhaps 70/30 or 80/20
- Cross-validation gives a way to use more of the data while still seeing how the model does on test data
  - Commonly 5 fold or 10 fold is done
  - Once a best model is chosen, model is refit on entire data set
- Often use both! Let's see why by introducing a model with a **tuning parameter**

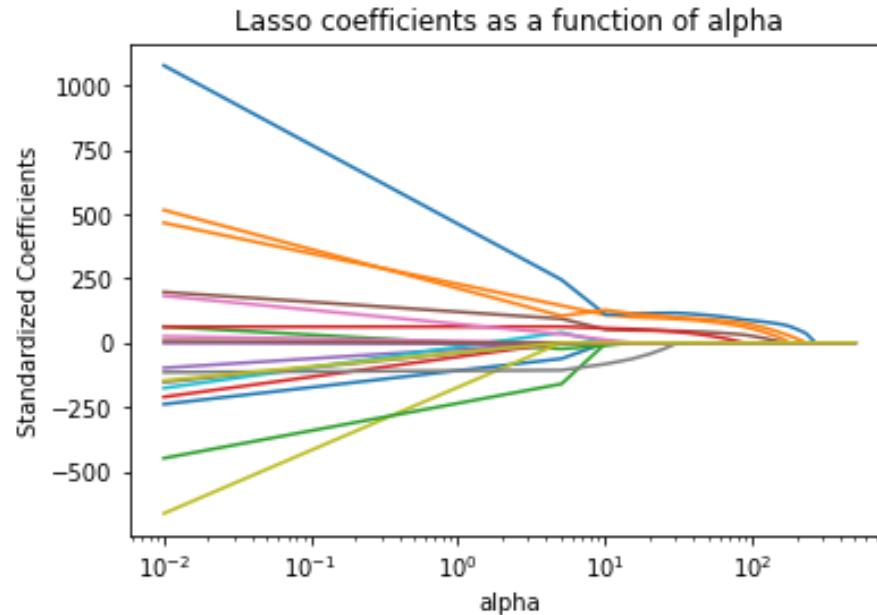
# LASSO Model

- Least Angle Subset and Selection Operator or LASSO
  - Similar to Least Squares but a penalty is placed on the sum of the absolute values of the regression coefficients
  - $\alpha (>0)$  is called a tuning parameter

$$\min_{\beta' s} \sum_{i=1}^n (y_i - (\beta_0 + \beta_1 x_{1i} + \dots + \beta_p x_{pi}))^2 + \alpha \sum_{j=1}^p |\beta_j|$$

# LASSO Model

- Least Angle Subset and Selection Operator or LASSO
  - Similar to Least Squares but a penalty is placed on the sum of the absolute values of the regression coefficients
  - Sets coefficients to 0 as you 'shrink'!



# Tuning Parameter

- When choosing the tuning parameter, we are really considering a **family of models!**
- Consider an  $\alpha = 0.1$  (small amount of shrinkage here)

```
from sklearn import linear_model
lasso = linear_model.Lasso(alpha=0.1)
lasso.fit(bike_data[["year", "log_km_driven"]].values, bike_data["log_selling_price"].values)

print(lasso.intercept_, lasso.coef_)

## -164.61209472866094 [ 0.08761607 -0.11092474]
```

# Tuning Parameter

- When choosing the tuning parameter, we are really considering a **family of models!**
- Consider an  $\alpha = 0.1$  (small amount of shrinkage here)

```
from sklearn import linear_model
lasso = linear_model.Lasso(alpha=0.1)
lasso.fit(bike_data[["year", "log_km_driven"]].values, bike_data["log_selling_price"].values)

print(lasso.intercept_, lasso.coef_)

## -164.61209472866094 [ 0.08761607 -0.11092474]
```

- Consider an  $\alpha = 1.05$  (a larger amount of shrinkage)

```
lasso = linear_model.Lasso(alpha=1.05)
lasso.fit(bike_data[["year", "log_km_driven"]].values, bike_data["log_selling_price"].values)

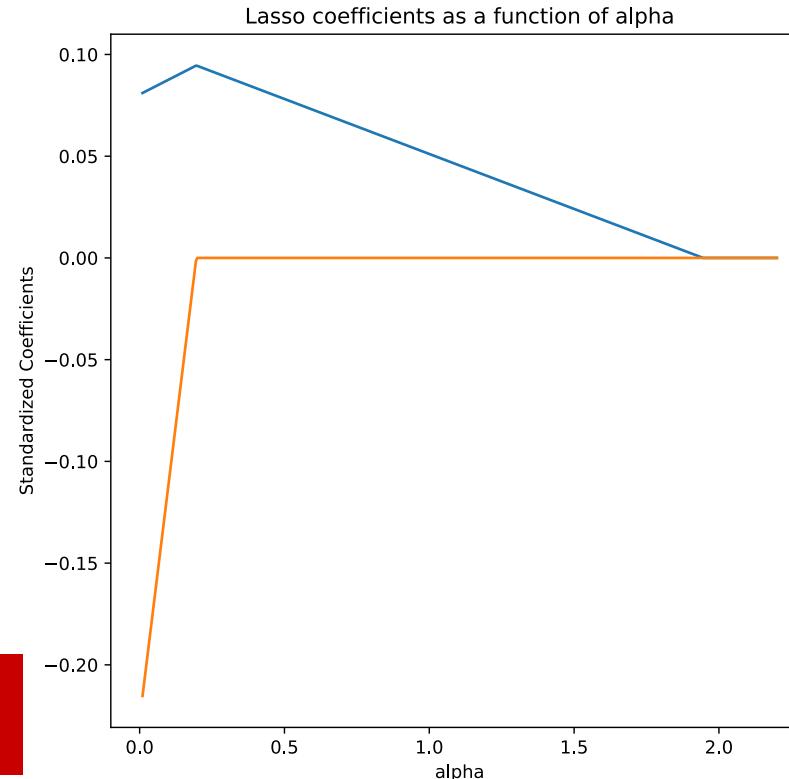
print(lasso.intercept_, lasso.coef_)

## -86.65630892150762 [ 0.04835598 -0. ]
```

# LASSO Fits Visual

- Perfect place for CV to help select the best  $\alpha$ !

```
## (-0.0995000000000003, 2.3095000000000003, -0.23074900910594773, 0.1099672686377952)
```



# Using CV to Select the Tuning Parameter

- Return the optimal  $\alpha$  using `LassoCV`

```
from sklearn.linear_model import LassoCV
lasso_mod = LassoCV(cv=5, random_state=0, alphas = np.linspace(0,2.2,100)) \
    .fit(bike_data[["year", "log_km_driven"]].values,
        bike_data["log_selling_price"].values)
```

# Using CV to Select the Tuning Parameter

- Return the optimal  $\alpha$  using `LassoCV`

```
pd.DataFrame(zip(lasso_mod.alphas_, lasso_mod.mse_path_), columns = ["alpha_value", "MSE_by_fold"])

##      alpha_value          MSE_by_fold
## 0    0.000000 [0.5496710875578059, 0.6805679103740427, 0.500...
## 1    0.022222 [0.5496710875578059, 0.6805679103740427, 0.500...
## 2    0.044444 [0.5496710875578059, 0.6805679103740427, 0.500...
## 3    0.066667 [0.5496710875578059, 0.6805679103740427, 0.500...
## 4    0.088889 [0.5496710875578059, 0.6805679103740427, 0.500...
## ..
## 95   2.111111 [0.3046546135682859, 0.3626276356362568, 0.191...
## 96   2.133333 [0.2998496943347488, 0.35411026477928204, 0.18...
## 97   2.155556 [0.29626758731408676, 0.34645956641174064, 0.1...
## 98   2.177778 [0.2939082925063055, 0.3396755405336282, 0.182...
## 99   2.200000 [0.29277214247542543, 0.33375857421410476, 0.1...

## [100 rows x 2 columns]
```

# Using CV to Select the Tuning Parameter

- Return the optimal  $\alpha$  using LassoCV

```
fit_info = np.array(list(zip(lasso_mod.alphas_, np.mean(lasso_mod.mse_path_, axis = 1))))  
fit_info[fit_info[:,0].argsort()]
```

```
## array([[0.          ,  0.26832555],  
##        [0.02222222,  0.26904464],  
##        [0.04444444,  0.27086204],  
##        [0.06666667,  0.27377741],  
##        [0.08888889,  0.27779157],  
##        [0.11111111,  0.28290414],  
##        [0.13333333,  0.28911508],  
##        [0.15555556,  0.29642441],  
##        [0.17777778,  0.30483239],  
##        [0.2        ,  0.30967893],  
##        [0.22222222,  0.31098715],  
##        [0.24444444,  0.31159763],  
##        [0.26666667,  0.31226415],  
##        [0.28888889,  0.31298674],  
##        [0.31111111,  0.31376537],  
##        [0.33333333,  0.31460007],  
##        [0.35555556,  0.31549081],  
##        [0.37777778,  0.31643761],  
##        [0.4        ,  0.31744046],  
##        [0.44444444,  0.31833333],  
##        [0.48888889,  0.32201242],  
##        [0.53333333,  0.32566667],  
##        [0.57777778,  0.32922222],  
##        [0.62222222,  0.33277778],  
##        [0.66666667,  0.33623333],  
##        [0.71111111,  0.33968889],  
##        [0.75555556,  0.34314444],  
##        [0.8        ,  0.34659999],  
##        [0.84444444,  0.35005556],  
##        [0.88888889,  0.35351111],  
##        [0.93333333,  0.35696667],  
##        [0.97777778,  0.36042222],  
##        [1.0        ,  0.36387778]]
```

# Using CV to Select the Tuning Parameter

- Now fit using that optimal  $\alpha$

```
lasso_best = linear_model.Lasso(lasso_mod.alpha_) #warning thrown since we are using 0, but can ignore
lasso_best.fit(bike_data[["year", "log_km_driven"]].values, bike_data["log_selling_price"].values)

print(lasso_best.intercept_, lasso_best.coef_)

## -148.7932910778814 [ 0.0803366 -0.22686129]
```

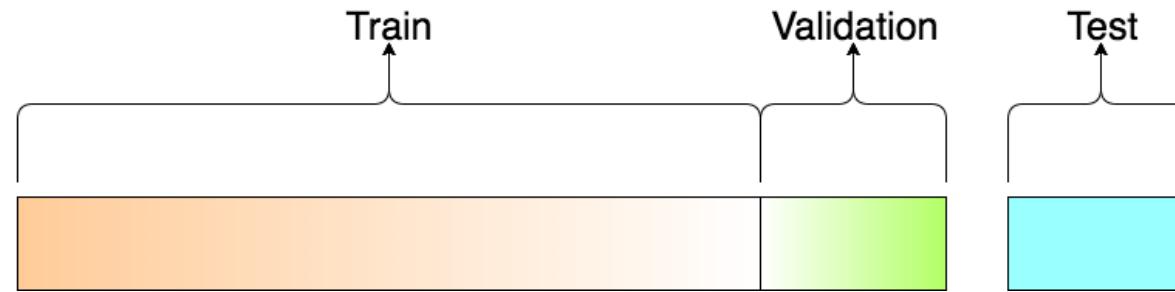
# So Do We Just Need CV?

Sometimes!

- If you are only considering one type of model, you can use just a training/test set or k-fold CV to select the best version of that model
- When you have multiple types of models to choose from, usually use both!
  - When we use the test set too much, we may have '**data leakage**'
  - Essentially we end up training our models to the test set by using it too much

# Training/Validation/Test or CV/Test

- Instead, we sometimes split into a training, validation, and test set
- CV can be used to replace the validation set!



# Training/Validation/Test or CV/Test

- Instead, we sometimes split into a training, validation, and test set
- CV can be used to replace the validation set!



- Compare only the **best** model from each model type on the test set

# Recap

- LASSO is similar to an MLR model but shrinks coefficients and may set some to 0
  - Tuning parameter must be chosen (usually by CV)
- Training/Test split gives us a way to validate our model's performance
  - CV can be used on the training set to select **tuning parameters**
  - Helps determine the 'best' model for a class of models
- With many competing model types, determine the best from each type check performance on the test set